

Conference writing and visualization plan

This extension turns the completed research release into a writing package. The experiment matrices, checkpoint choices, frozen protocols and reported scores remain those of release v1.0.0. No additional model training, inference, test-set tuning or hypothesis selection is part of this extension. A separate Kaggle CPU notebook renders the figure collection from the frozen evidence.

The venue and page limit have not been specified. Use a venue-neutral manuscript structure and select a compact main-paper figure set; retain the full 20-figure collection as supporting material. The principal study concerns missing-modality segmentation. The small classification benchmark is separate supporting work and must not be merged into one headline performance claim.

Figure	Evidence and purpose	Initial writing placement
01	Dataset counts before and after similarity grouping	Data audit
02	Executed fold-0 roles and reserved cohort	Evaluation design
03	All seven classification variants	Classification supplement
04	Per-seed OOF ROC and precision–recall curves	Classification supplement
05	Three confusion matrices at the frozen threshold	Classification supplement
06	All six paired classifier contrasts	Classification supplement
07	All nine main segmentation variants and three endpoints	Main ablation study
08	Adjusted intervals for all main QMMF-control contrasts	Main ablation study
09	All 27 inner-selection learning curves	Optimization supplement
10	Capacity, score and recorded compute	Resource supplement
11	All four reserved model families and three endpoints	Primary results
12	All declared reserved paired contrasts	Primary comparisons
13	WT, TC and ET performance with the empty-reference rule	Regional results
14	All 15 modality combinations for all four models	Robustness
15	Modality Shapley values under the declared empty utility	Exploratory attribution
16	All 51 paired group scores and their differences	Error distribution
17	Both ET-empty cases, all models and seeds	Failure analysis
18	Complete cohort and frozen historical sensitivity	Sensitivity / limitations
19	Unmodified prespecified real prediction montage	Qualitative results
20	Measured lossless-cache throughput	Implementation supplement

Each figure receives PNG, PDF and SVG files, a plotted-data CSV, a full caption, a limited interpretation and source-file hashes. Figure 19 preserves the original PNG/PDF bytes and attribution; its SVG is a wrapper around the original raster. Other chart PDF/SVG exports preserve vector elements. The executed Kaggle gallery is retained separately from the portable notebook source.

The paper guide and draft must retain the observed quality-module failure, unknown patient identity, historical validation exposure, simulated missingness after complete-modality preprocessing, fixed training seeds, within-endpoint multiplicity adjustment, auxiliary-objective differences between model families, and lack of independent external validation. Figure count is not a measure of scientific novelty or conference acceptance.

Validation covers source hashes, plotted values, complete seeds/models/groups, image and PDF rendering, readable labels, real Kaggle execution, notebook output images, relative GitHub links and downloaded release checksums. Publication uses a new versioned GitHub release and preserves v1.0.0.